

# Kubernetes deployment with Terraform on cloud.ca

Clément Contini

January 26, 2018



# Objectives

**Objective:** Present how to use an *infrastructure as code* tool to deploy a Kubernetes cluster.

This presentation will consist of:

- ▶ Explaining the technologies involved
- ▶ Detailing the Terraform configuration used
- ▶ Seeing the results in a demo

# About us

## *Clément Contini*

- ▶ Software Engineering degree from ÉTS
- ▶ I work on Cloud Infrastructure at CloudOps and cloud.ca

## *CloudOps*

- ▶ Cloud Infrastructure experts since 2005
- ▶ Design, build and manage public, private and hybrid cloud solutions

## *cloud.ca*

- ▶ Infrastructure as a service platform, based in Montreal
- ▶ Launched in August 2014

# Terraform

Infrastructure as code tool:

- ▶ Cross-platform
- ▶ Configurations can be versioned and reused
- ▶ Generates execution plans
- ▶ A cloud.ca provider is available

Similar tools:

- ▶ OpenStack Heat
- ▶ AWS CloudFormation



**Figure 1:** Terraform logo

# Kubernetes

- ▶ Container orchestration system
- ▶ Originates from Google's Borg

Similar tools:

- ▶ Docker Swarm
- ▶ ~ Rancher
- ▶ ~ HashiCorp Nomad



**Figure 2:** Kubernetes logo

# Demonstration

The demo will consist in deploying a Kubernetes cluster using the *kubeadm* toolkit.

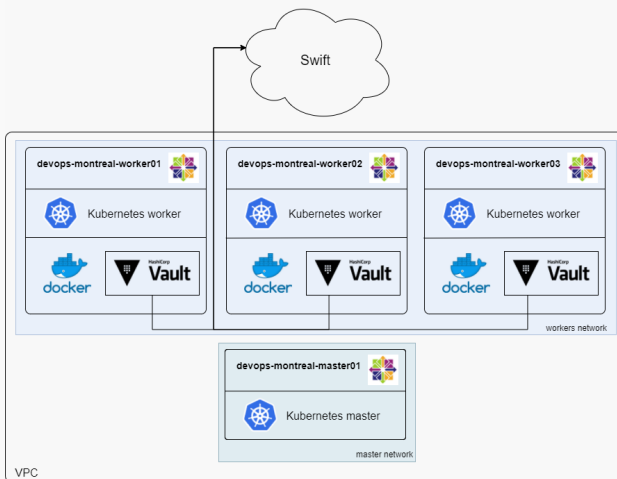
We will deploy:

- ▶ a cloud.ca environment with a VPC, networks and all the necessary virtual machines
- ▶ A Kubernetes master and several workers on the VMs
- ▶ Automatically start a Vault deployment in the cluster with 3 replicas

Let's start the demo, and I will explain the configuration while it's building the cluster.



# Demonstration



**Figure 3:** Deployment of the Kubernetes cluster

# Environment

We first create an environment in cloud.ca for our deployment.

We create:

- ▶ a cloudca\_environment
- ▶ a cloudca\_vpc
- ▶ 2 cloudca\_public\_ip

```
1 resource "cloudca_environment" "kubernetes" {  
2     service_code      = "${var.service_code}"  
3     organization_code = "${var.organization_code}"  
4     name              = "${format("%s-env", var.prefix)}"  
5     description       = "Environment for a Kubernetes cluster"  
6     admin_role        = ["${var.admin}"]  
7     read_only_role    = ["${var.read_only}"]  
8 }
```

## Listing 1: Environment definition



# Variables

Some variables need to be defined (for example the api key) and some can simply use the default. The values are provided in the *terraform.tfvars* file.

```
1 | service_code      = "compute-qc"  
2 | organization_code = "system"  
3 | admin             = [ "ccontini" ]  
4 | read_only         = [ ]  
5 | prefix            = "devops-montreal"
```

**Listing 2:** Variables assignment

# Variables

We need to declare the variables that will help us configure the environment for our use case:

- ▶ service\_code
- ▶ organization\_code
- ▶ prefix
- ▶ admin
- ▶ read\_only

```
1 variable "service_code" {}  
2 variable "organization_code" {}  
3 variable "admin" { type = "list" }  
4 variable "read_only" { type = "list" }  
5 variable "prefix" {}
```

**Listing 3:** Variables definition

# Master

We then declare all resources required to start the Kubernetes master:

- ▶ A network and its associated set of ACLs
- ▶ An instance
- ▶ A port forwarding rule

```
1 resource "cloudca_network" "master" {  
2   environment_id = "${cloudca_environment.kubernetes.id}"  
3   name           = "${format("%s-network-master", var.  
    ↪ prefix)}"  
4   description    = "Network for the master nodes"  
5   vpc_id         = "${cloudca_vpc.kubernetes.id}"  
6   network_offering = "Standard Tier"  
7   network_acl_id  = "${cloudca_network_acl.master.id}"  
8 }
```

**Listing 4:** Network definition for the master node

# Master

```
1 resource "cloudca_instance" "master_node" {
2     environment_id      = "${cloudca_environment.kubernetes
      ↪     .id}"
3     name                = "${format("%s-master01", var.
      ↪     prefix)}"
4     network_id          = "${cloudca_network.master.id}"
5     template            = "${var.template_name}"
6     compute_offering    = "${var.default_offering}"
7     cpu_count           = "${var.master_vcpu}"
8     memory_in_mb        = "${var.master_ram}"
9     root_volume_size_in_gb = "${var.master_disk}"
10    user_data            = "${data.template_file.cloudinit.
      ↪    rendered}"
11 }
```

**Listing 5:** Master node instance definition

# Master

The master node is configured using a cloudinit script. It will:

- ▶ Create a user
- ▶ Install docker and Kubernetes
- ▶ start a bootstrap script to initialize the master

```
1 data "template_file" "cloudinit" {  
2   template = "${file("templates/cloudinit.tpl")}"  
3  
4   vars {  
5     public_key = "${replace(file("./id_rsa.pub"), "\n", "")}"  
6     username   = "${var.username}"  
7   }  
8 }
```

**Listing 6:** Cloudinit template interpolation

# Master

```
1 resource "cloudca_port_forwarding_rule" "  
    ↪ management_master_ssh" {  
2     environment_id      = "${cloudca_environment.kubernetes.id}  
    ↪ "  
3     public_ip_id        = "${cloudca_public_ip.master_ip.id}"  
4     public_port_start   = 2200  
5     private_ip_id       = "${cloudca_instance.master_node.  
    ↪ private_ip_id}"  
6     private_port_start  = 22  
7     protocol            = "TCP"
```

**Listing 7:** Master node port forwarding rule

Once the port forwarding rule is created, we use it to provision the node with bootstrap scripts and a deployment manifest that will be applied when the Kubernetes cluster starts.

# Master

```
1  connection {
2      type      = "ssh"
3      user      = "${var.username}"
4      private_key = "${file("./id_rsa")}"
5      host      = "${cloudca_public_ip.master_ip.ip_address}
6      ↪      "
7      port      = 2200
8  }
9
10  provisioner "file" {
11      content      = "${data.template_file.bootstrap_master.
12      ↪      rendered}"
13      destination = "/home/${var.username}/bootstrap.sh"
14  }
15
16  provisioner "file" {
17      source      = "manifests"
18      destination = "/home/${var.username}/manifests"
```

**Listing 8:** Provisioning of the master node

# Workers

We create the workers, the same way we created the master, with:

- ▶ A network and its associated set of ACLs
- ▶ Several instance (we will use the *count* keyword)
- ▶ A port forwarding rule

```
1 resource "cloudca_instance" "worker_nodes" {  
2   environment_id      = "${cloudca_environment.kubernetes  
   ↪ .id}"  
3   name                = "${format("%s-worker%02d", var.  
   ↪ prefix, count.index + 1)}"  
4   network_id         = "${cloudca_network.worker.id}"  
5   template            = "${var.template_name}"  
6   compute_offering    = "${var.default_offering}"  
7   cpu_count          = "${var.worker_vcpu}"  
8   memory_in_mb       = "${var.worker_ram}"  
9   root_volume_size_in_gb = "${var.worker_disk}"  
10  count               = "${var.worker_count}"  
11  user_data            = "${data.template_file.cloudinit.  
   ↪ rendered}"  
12 }
```

**Listing 9:** Workers instance definition



# Vault deployment

Let's check what was deployed on Kubernetes.

## Some alternatives

Some alternative solutions that could be use to create and manage a Kubernetes cluster:

- ▶ We could use the Kubernetes Terraform provider to manage the resources inside the Kubernetes cluster.
- ▶ We could use Rancher to deploy Kubernetes on cloud.ca using a GUI

*cloud.ca* plans to integrate more tightly with Kubernetes to be able to create resources for the containers directly with Kubernetes:

- ▶ Volumes
- ▶ Load-balancing rules

# Questions

Any questions ?

Thanks for attending this meetup !

The source code is available at:

<https://github.com/vilisseranen/terraform-kubeadm>.

